

Beyond Convolutions: A Deep Dive into Vision Transformer (ViT) Architectures, Mathematical Foundations, and Scaling Comparative Analysis

Project Lead: Ahmed Farouk
Faculty of Computer Science & Engineering
Alamein University

April 19, 2026

Abstract

The Vision Transformer (ViT) represents a paradigm shift in computer vision, moving away from the local inductive biases of Convolutional Neural Networks (CNNs) toward a global self-attention mechanism. This paper provides a rigorous mathematical definition of the ViT architecture, exploring the mechanics of patch embedding and multi-headed self-attention. We conduct a comparative study across three scaling variants: ViT-Base, ViT-Large, and ViT-Huge, evaluating their performance metrics, computational trade-offs, and data dependencies. Finally, we discuss the inherent drawbacks of the transformer paradigm and propose a vision for future enhancements, specifically focusing on hybrid architectures and efficient attention mechanisms.

1 Introduction

Since the inception of AlexNet, Convolutional Neural Networks (CNNs) have dominated the field of computer vision. However, the introduction of the Vision Transformer (ViT) by Dosovitskiy et al. (2020) demonstrated that the Transformer architecture—the backbone of modern Natural Language Processing (NLP)—could be applied directly to image data with remarkable success. By treating an image as a sequence of patches, ViTs leverage global dependencies that are often missed by the local receptive fields of traditional kernels.

2 The Vision Transformer Architecture

The ViT architecture is designed to process 2D images using a 1D sequence-to-sequence transformer encoder. The process involves four critical stages:

2.1 Image Pacification and Linear Projection

An input image $\mathbf{x} \in \mathbb{R}^{H \times W \times C}$ is decomposed into N non-overlapping patches $\mathbf{x}_p \in \mathbb{R}^{N \times (P^2 \cdot C)}$, where P is the patch resolution and $N = HW/P^2$. These patches are flattened and mapped to a D -dimensional latent space using a trainable linear projection matrix $\mathbf{E}$.

2.2 Position Embeddings and Class Token

To perform classification, a learnable $[class]$ token $\mathbf{x}_{class}$ is prepended to the patch sequence. Because Transformers are permutation-invariant, 1D learnable positional embeddings $\mathbf{E}_{pos}$ are added to retain spatial information:

$$\mathbf{z}_0 = [\mathbf{x}_{class}; \mathbf{x}_p^1 \mathbf{E}; \dots; \mathbf{x}_p^N \mathbf{E}] + \mathbf{E}_{pos} \quad (1)$$

2.3 The Transformer Encoder

The encoder consists of L identical layers. Each layer contains two sub-blocks: Multi-Headed Self-Attention (MSA) and a Multilayer Perceptron (MLP). Layer Normalization (LN) is applied before each sub-block, and residual connections are applied after.

3 Mathematical Foundations

The robustness of ViTs stems from the self-attention mechanism, which allows every patch to interact with every other patch in the sequence.

3.1 Multi-Headed Self-Attention (MSA)

For each head i , we compute Query ($\mathbf{Q}$), Key ($\mathbf{K}$), and Value ($\mathbf{V}$) matrices:

$$\mathbf{Q}_i = \mathbf{z} \mathbf{W}_i^Q, \quad \mathbf{K}_i = \mathbf{z} \mathbf{W}_i^K, \quad \mathbf{V}_i = \mathbf{z} \mathbf{W}_i^V \quad (2)$$

The attention weights are calculated using the scaled dot-product:

$$\text{Attention}(\mathbf{Q}_i, \mathbf{K}_i, \mathbf{V}_i) = \text{softmax}\left(\frac{\mathbf{Q}_i \mathbf{K}_i^T}{\sqrt{d_k}}\right) \mathbf{V}_i \quad (3)$$

The MSA output is the concatenation of all heads projected through $\mathbf{W}^O$:

$$\text{MSA}(\mathbf{z}) = [\text{head}_1; \dots; \text{head}_k] \mathbf{W}^O \quad (4)$$

3.2 MLP and Layer Equations

The MLP consists of two linear transformations with a GELU activation:

$$\text{MLP}(\mathbf{z}) = \text{GELU}(\mathbf{z} \mathbf{W}_1 + \mathbf{b}_1) \mathbf{W}_2 + \mathbf{b}_2 \quad (5)$$

The full state transition for layer ℓ is:

$$\mathbf{z}'_\ell = \text{MSA}(\text{LN}(\mathbf{z}_{\ell-1})) + \mathbf{z}_{\ell-1} \quad (6)$$

$$\mathbf{z}_\ell = \text{MLP}(\text{LN}(\mathbf{z}'_\ell)) + \mathbf{z}'_\ell \quad (7)$$

4 Scaling and Comparative Analysis

We evaluate three primary variants of the Vision Transformer, differing in depth (L), hidden dimension (D), and the number of attention heads.

Table 1: ViT Model Variant Comparison

Model	Layers	D	Heads	Params
ViT-Base	12	768	12	86M
ViT-Large	24	1024	16	307M
ViT-Huge	32	1280	16	632M

4.1 Analysis of Performance

1. **ViT-Base:** Efficient for general tasks. Our experiments with CLIP-ViT-B/32 show high robustness but slightly lower granularity on complex scenes. 2. **ViT-Large:** Demonstrates significant gains in zero-shot accuracy. The reduced patch size (e.g., $L/14$) allows for superior semantic capture. 3. **ViT-Huge:** The frontier of vision modeling. It excels in large-scale pre-training (JFT-300M) but suffers from overfitting on smaller datasets without heavy regularization.

5 Drawbacks and Advantages

5.1 Advantages

- **Global Receptive Field:** Interaction between distant pixels from the first layer.

- **Multimodal Compatibility:** Identical structure to LLMs, enabling unified image-text embedding spaces.

5.2 Drawbacks

- **Data Hunger:** Requires massive datasets to learn spatial relationships.
- **Quadratic Complexity:** Memory usage scales as $O(N^2)$, making high-res imagery expensive.

6 Vision for Enhancement

To advance the state of ViTs, we propose: 1. **Convolutional Stemming:** Using 3x3 convolutions for the initial patch embedding to reintroduce local inductive bias. 2. **Linear Attention:** Replacing softmax attention with kernel-based linear approximations to reduce complexity to $O(N)$. 3. **Masked Image Modeling:** Leveraging MAE-style training to reduce reliance on labeled data.

7 Conclusion

The Vision Transformer has fundamentally altered the computer vision landscape. While challenges remain in data efficiency and computational cost, its scaling properties and multimodal flexibility make it the definitive architecture for the next generation of artificial intelligence.